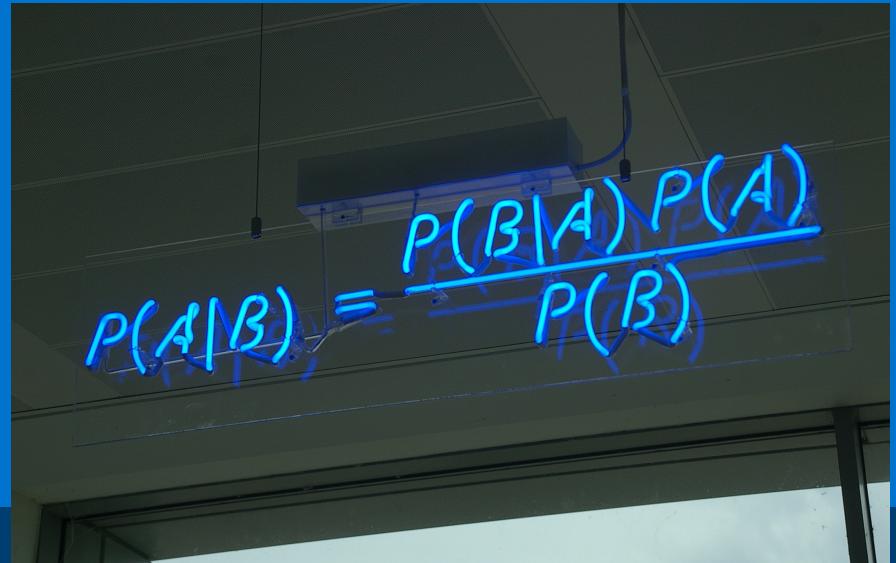


Bayesian Methods in Deep Learning

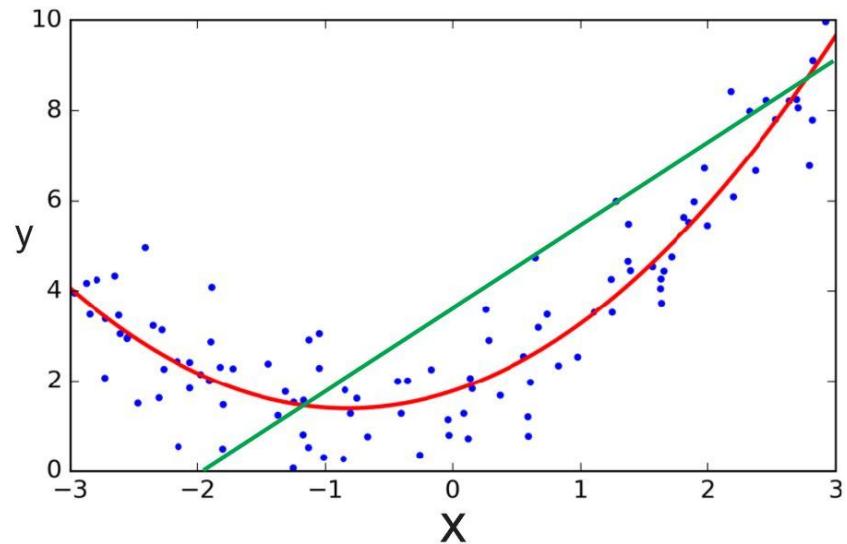
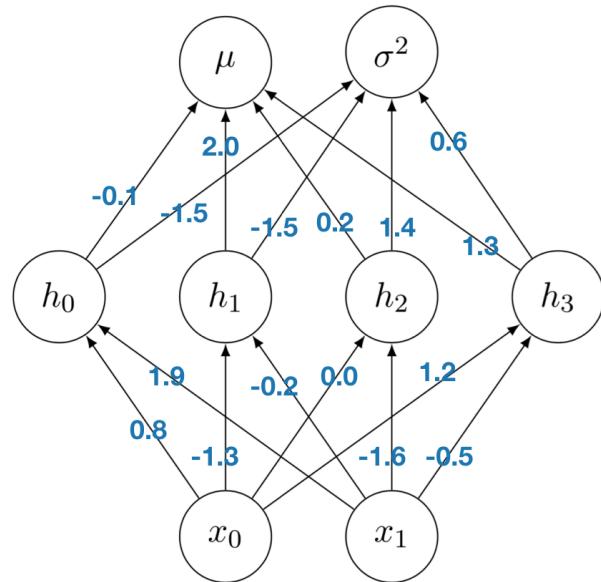
Javier Antorán (ja666@cam.ac.uk)

Feel free to interrupt at any time :)



Deep learning

- Overparametrised non-linear models
- ML / MAP inference + Stochastic optimisation
- Needs lots of data

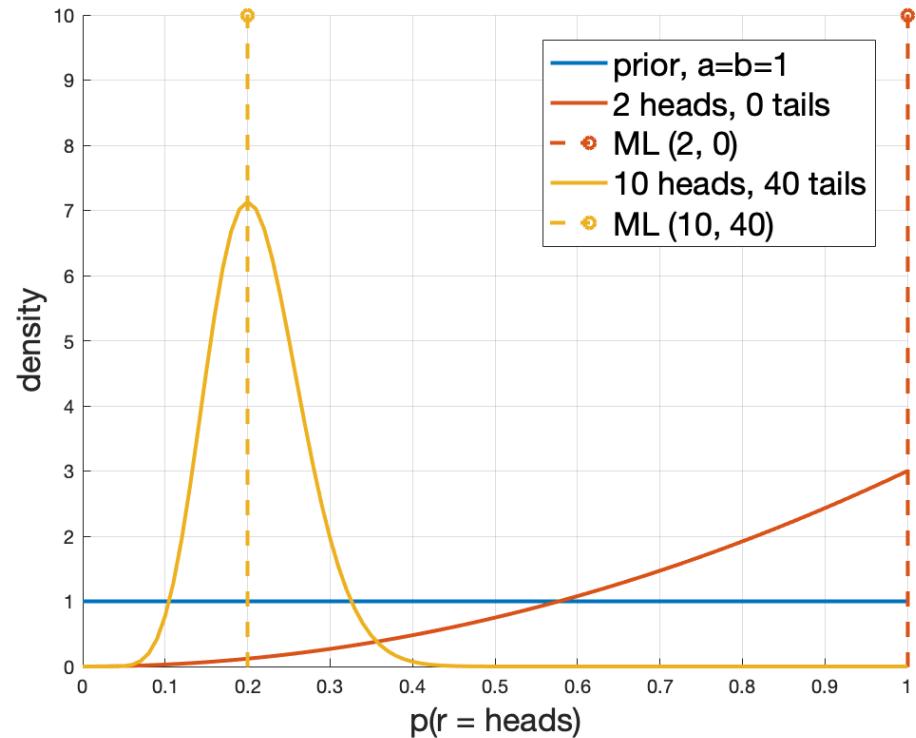


Probabilistic Inference: A biased coin

Likelihood

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})p(\mathbf{w})}{p(\mathcal{D})}$$

Prior

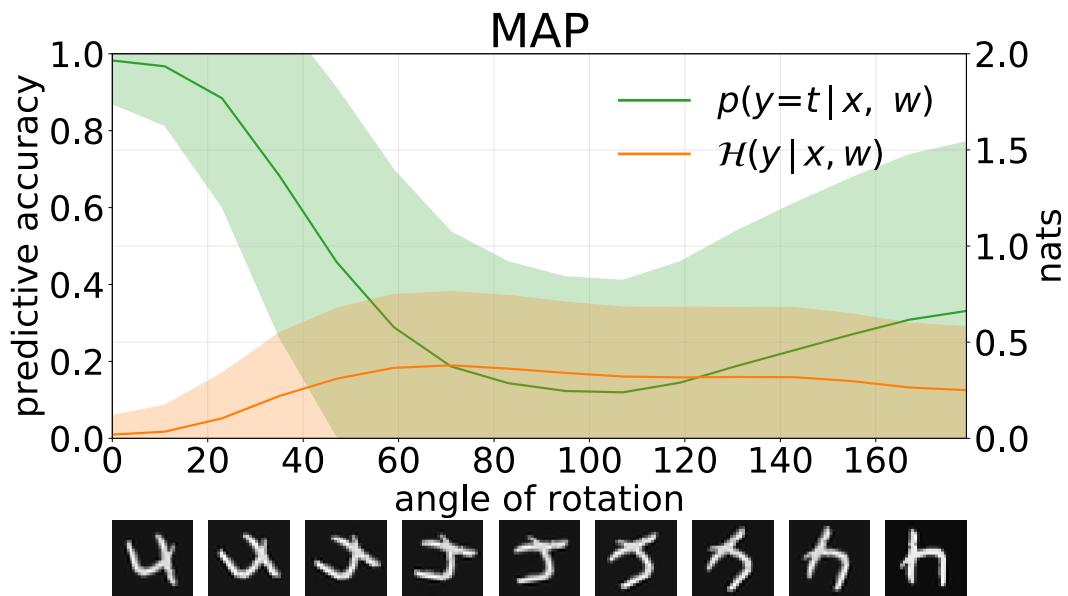


The Problem With MAP / ML

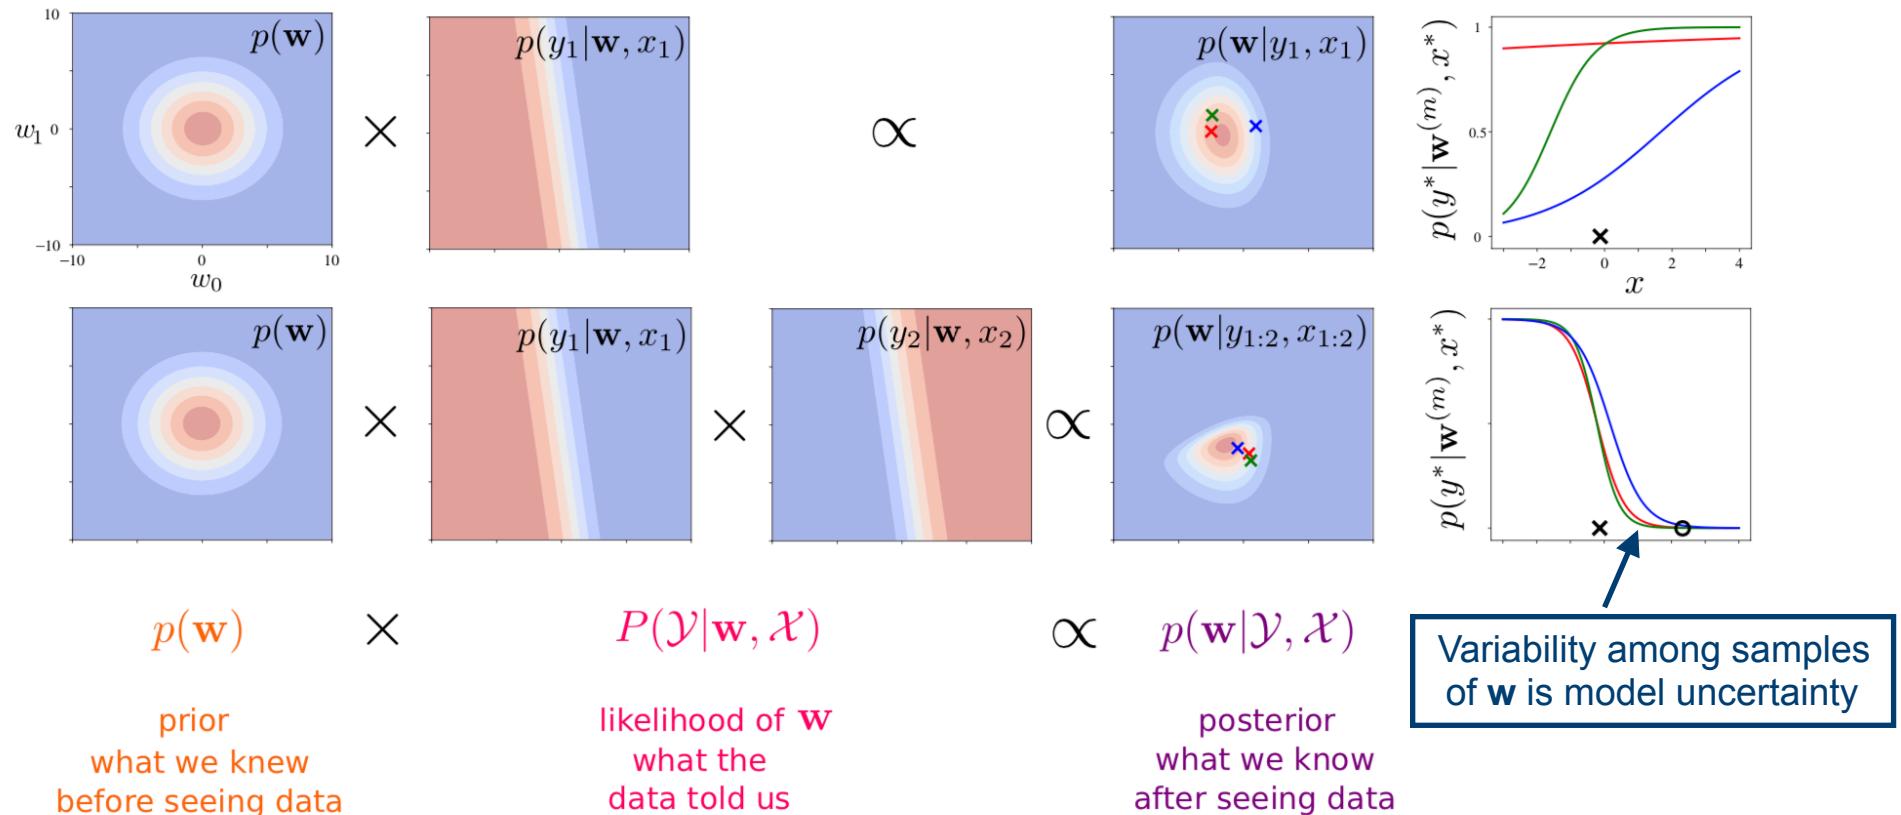
Given a dataset $\mathcal{D} = \{\mathbf{X}, \mathbf{Y}\}$ composed of inputs $\mathbf{X} = \{\mathbf{x}_1 \dots \mathbf{x}_n\}$ and targets $\mathbf{Y} = \{\mathbf{y}_1 \dots \mathbf{y}_n\}$, and a neural network parametrised by \mathbf{w} , the maximum likelihood criterion:

$$\mathbf{w}_{ml} = \arg \max_{\mathbf{w}} p(\mathcal{D} | \mathbf{w}) = \arg \max_{\mathbf{w}} \log p(\mathcal{D} | \mathbf{w}) \quad (2.1)$$

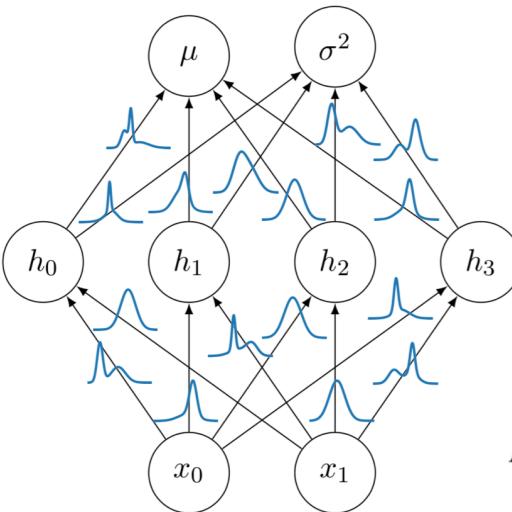
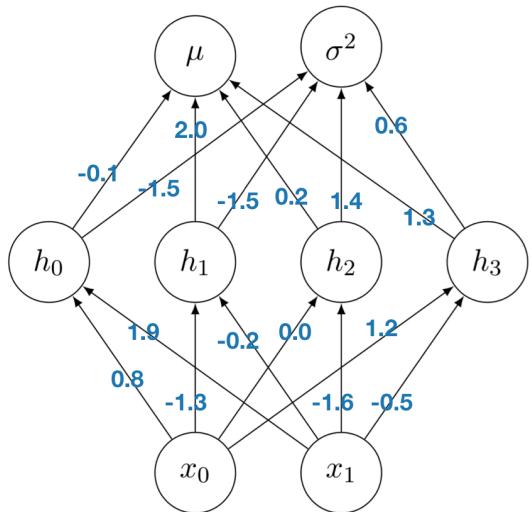
- MAP / ML returns the parameter setting that best explains the data.
- It may not be the only good explanation.
- No guarantees about unseen data. These must be introduced through priors.



1-d Probabilistic Logistic Regression

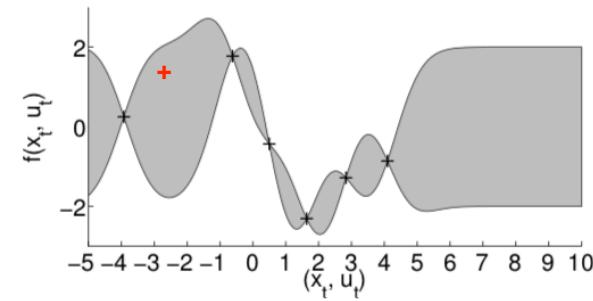
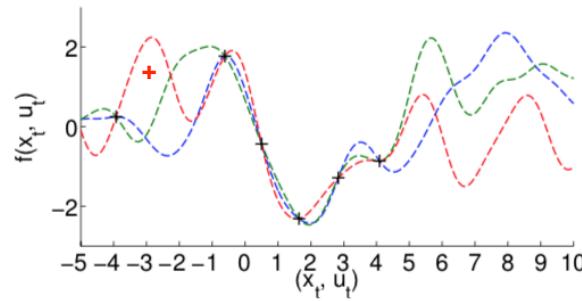
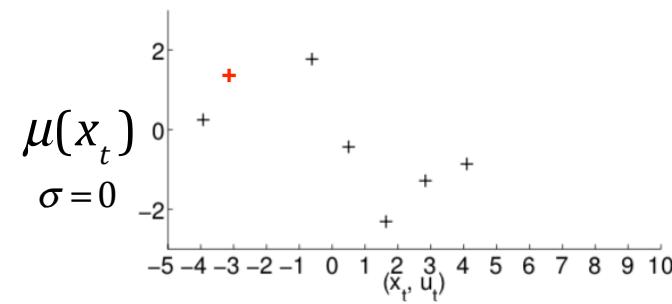


Bayesian Neural Networks (BNNs)



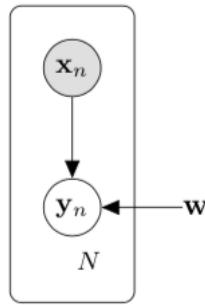
- Predictions incorporate model uncertainty!

$$p(\mathbf{Y}^* | \mathbf{X}^*, \mathcal{D}) = \int p(\mathbf{Y}^* | \mathbf{X}^*, \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w}$$

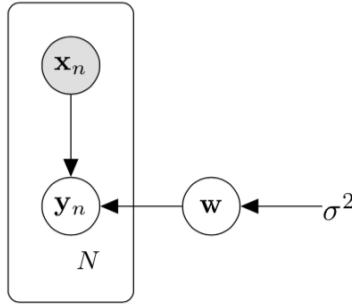


Practical Forms of Prior

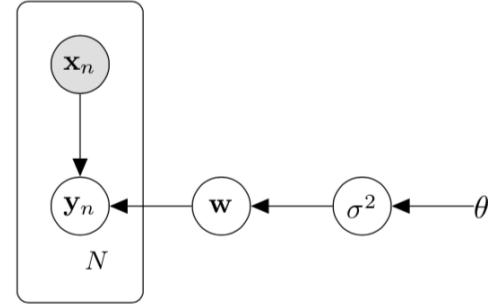
- ML / MAP a)
- Symmetric tractable prior b)
 - Gaussian, Laplace, Student-T, etc
 - Type-2 Maximum Likelihood
 - Previous posterior (Bayesian update)
- Conjugate hyper-prior + Gibbs sampling (hierarchical model) c)



a)



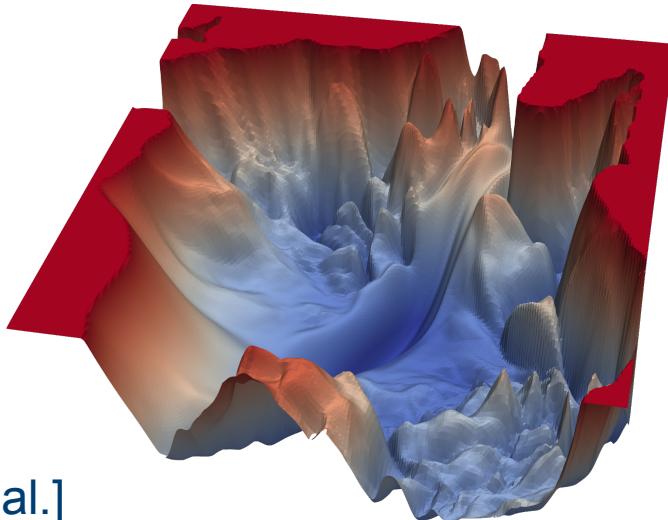
b)



c)

Inference with BNNs

- The data's likelihood under a BNN model is a very complex, high dimensional and multimodal function.



[Li et al.]

- **Intractable evidence:**

$$p(\mathcal{D}) = \int p(\mathcal{D}|\mathbf{w})p(\mathbf{w}) d\mathbf{w}$$

- **Intractable predictive:**

$$p(\mathbf{Y}^*|\mathbf{X}^*, \mathcal{D}) = \int p(\mathbf{Y}^*|\mathbf{X}^*, \mathbf{w})p(\mathbf{w}|\mathcal{D}) d\mathbf{w}$$

- Must resort to **Approximate Inference**

Variational Inference

Calculus of variations is like finding optima of a function but instead of finding a value of (x, y) which yields the optima, you find a function which yields the optima.

Kullback–Leibler (KL) divergence

$$\text{KL}(q(\mathbf{w})||p(\mathbf{w})) = \int q(\mathbf{w}) \log \frac{q(\mathbf{w})}{p(\mathbf{w})} d\mathbf{w}$$

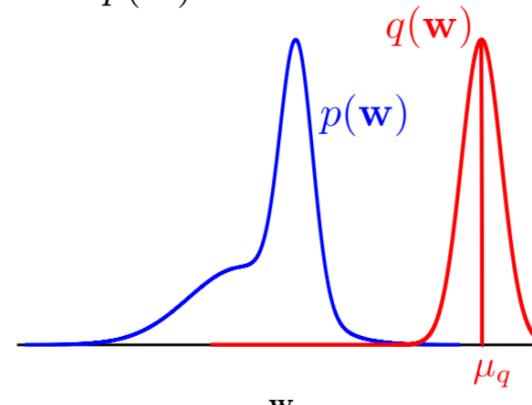
1. non-negative

$$\frac{\delta^2}{\delta q^2} \text{KL}(q(\mathbf{w})||p(\mathbf{w})) = \frac{1}{q(\mathbf{w})} \geq 0$$

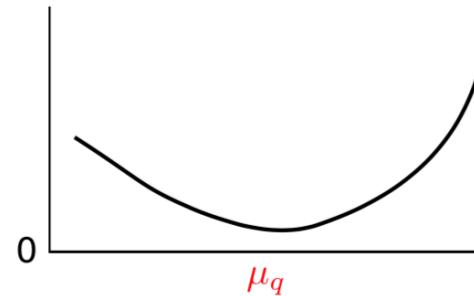
2. zero (minimised) when $q(\mathbf{w}) = p(\mathbf{w})$

$$\text{KL}(p(\mathbf{w})||p(\mathbf{w})) = \int p(\mathbf{w}) \log \frac{p(\mathbf{w})}{p(\mathbf{w})} d\mathbf{w} = 0$$

First grad = 0 only if $q=p$.
Second grad \rightarrow strictly convex



$$\text{KL}(q(\mathbf{w})||p(\mathbf{w}))$$

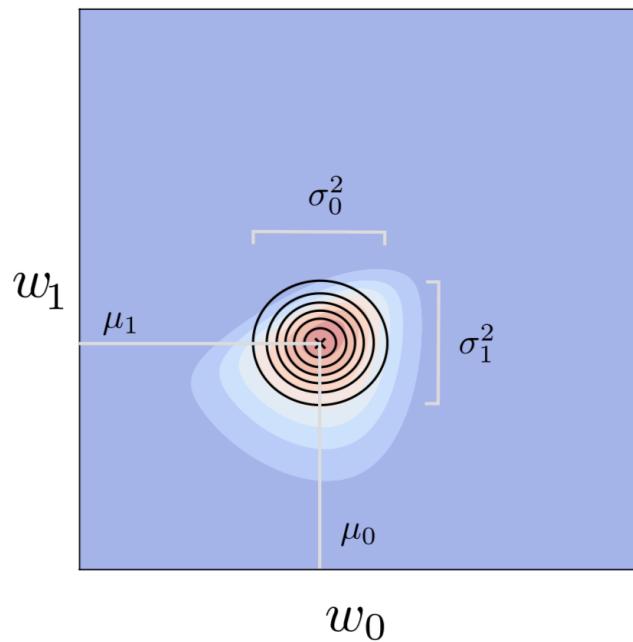


3. can be computed up to an additive constant w/o needing normalisation for $p(\mathbf{w})$

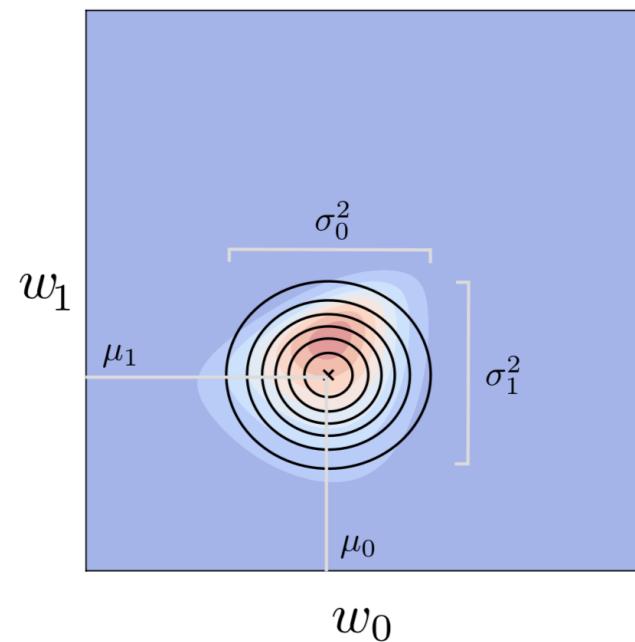
Issues with VI

- Mode seeking behaviour
- Overly simple approximate posteriors (complexity vs tractability)

Laplace Approximation (diagonal Hessian)



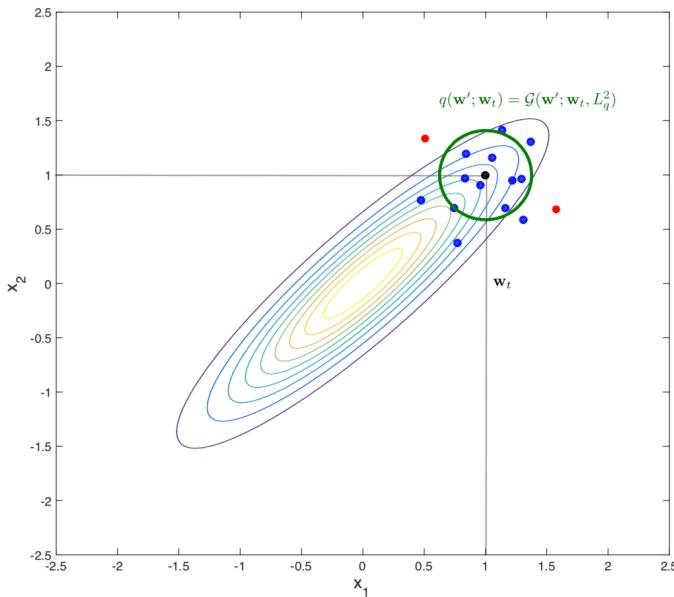
Variational Inference (factorised)



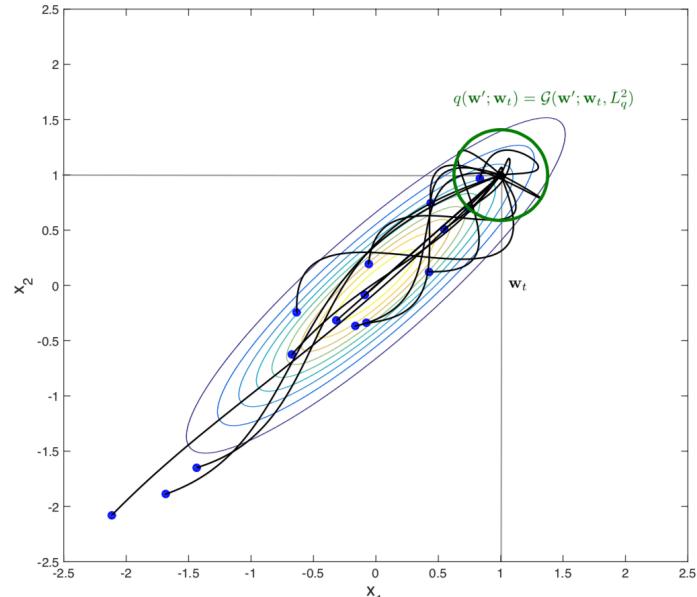
Markov Chain Monte Carlo

$$p(\mathbf{Y}^* | \mathbf{X}^*, \mathcal{D}) \approx \sum_{k=1}^K p(\mathbf{Y}^* | \mathbf{X}^*, \mathbf{w}_k); \quad \mathbf{w}_k \sim p(\mathbf{w} | \mathcal{D})$$

- Robust uncertainty requires sample diversity: [HarleMCMC shake](#)

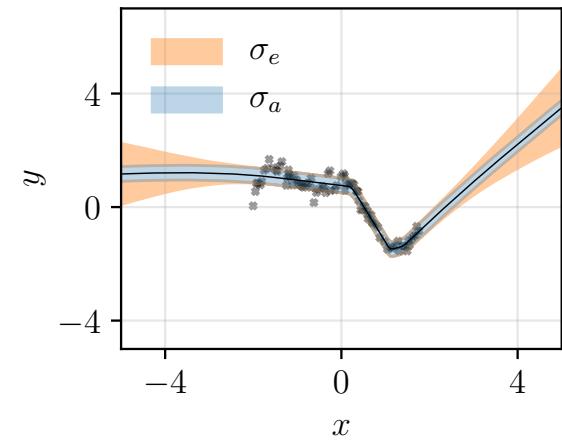
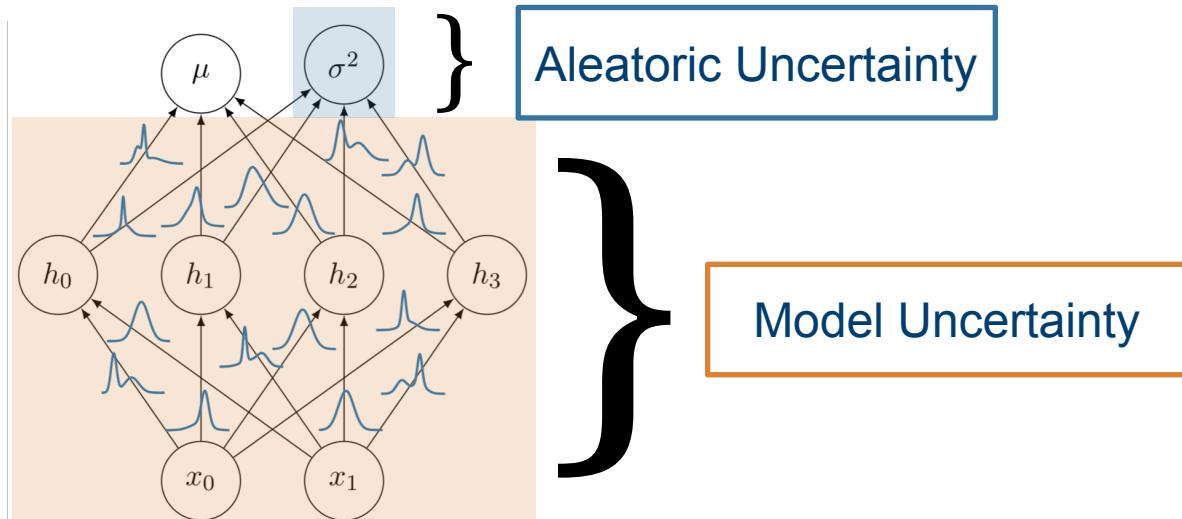


Metropolis

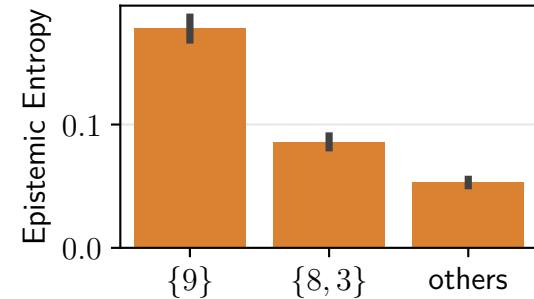
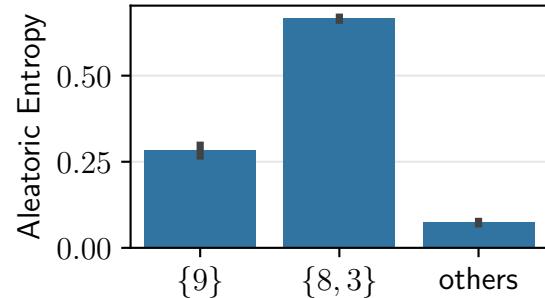


Hybrid Monte Carlo

Uncertainty Decomposition

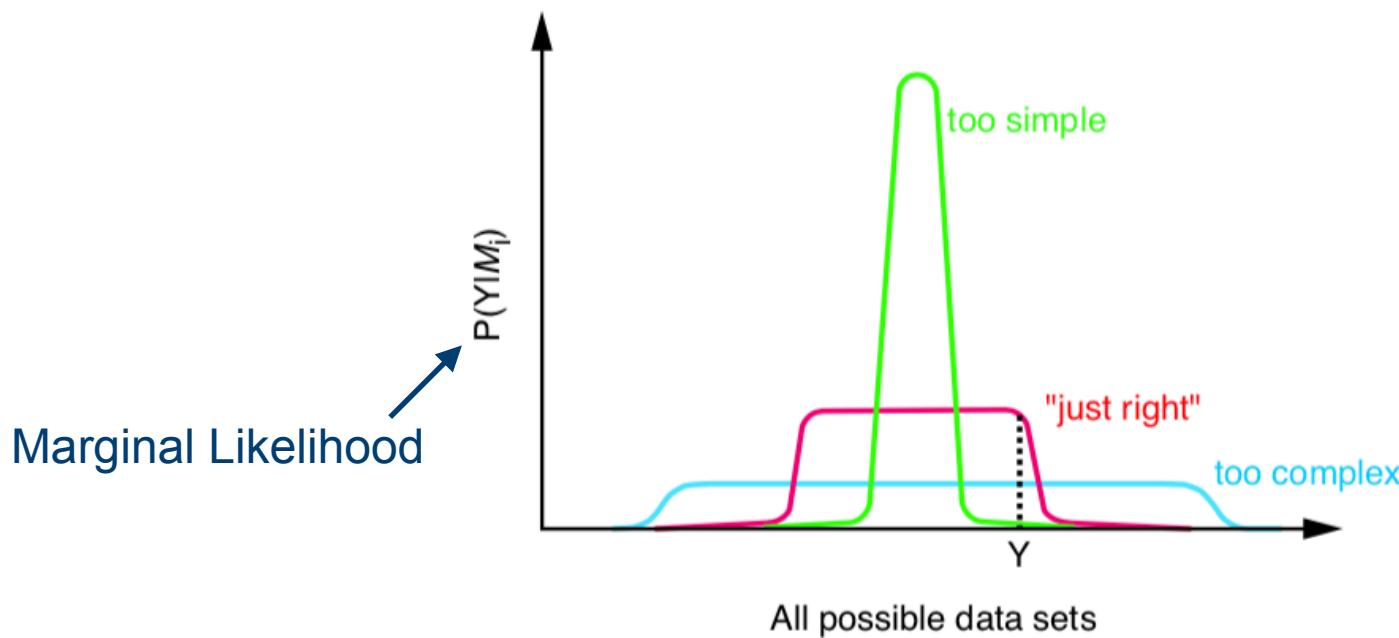


MNIST; Remove 80% of
{9} and permute 30% {8,3}



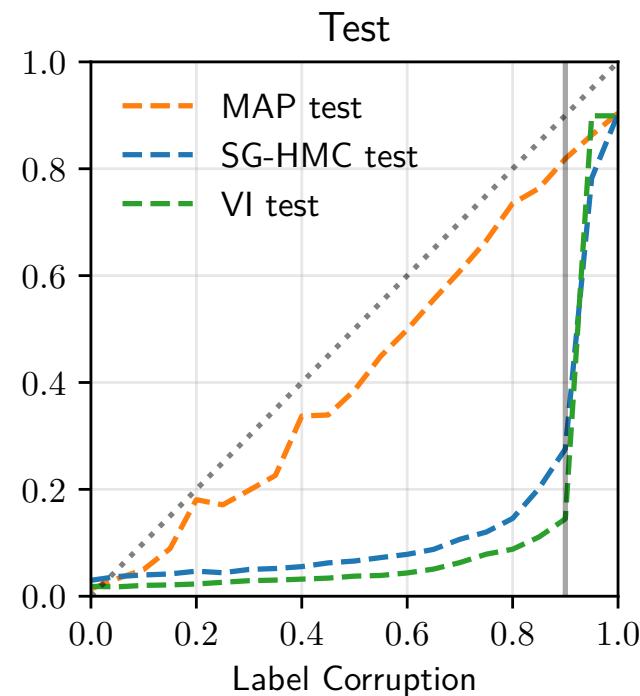
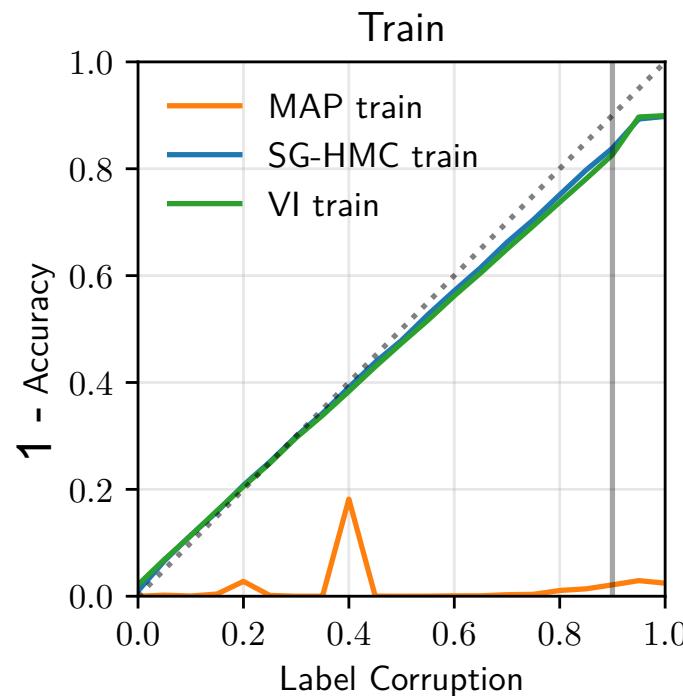
Protection Against Overfitting

- Model uncertainty is transformed into predictive uncertainty
- With uninformative prior, automatic tradeoff between goodness of fit and model complexity: Automatic Ockham's Razor. (Related to MDL)



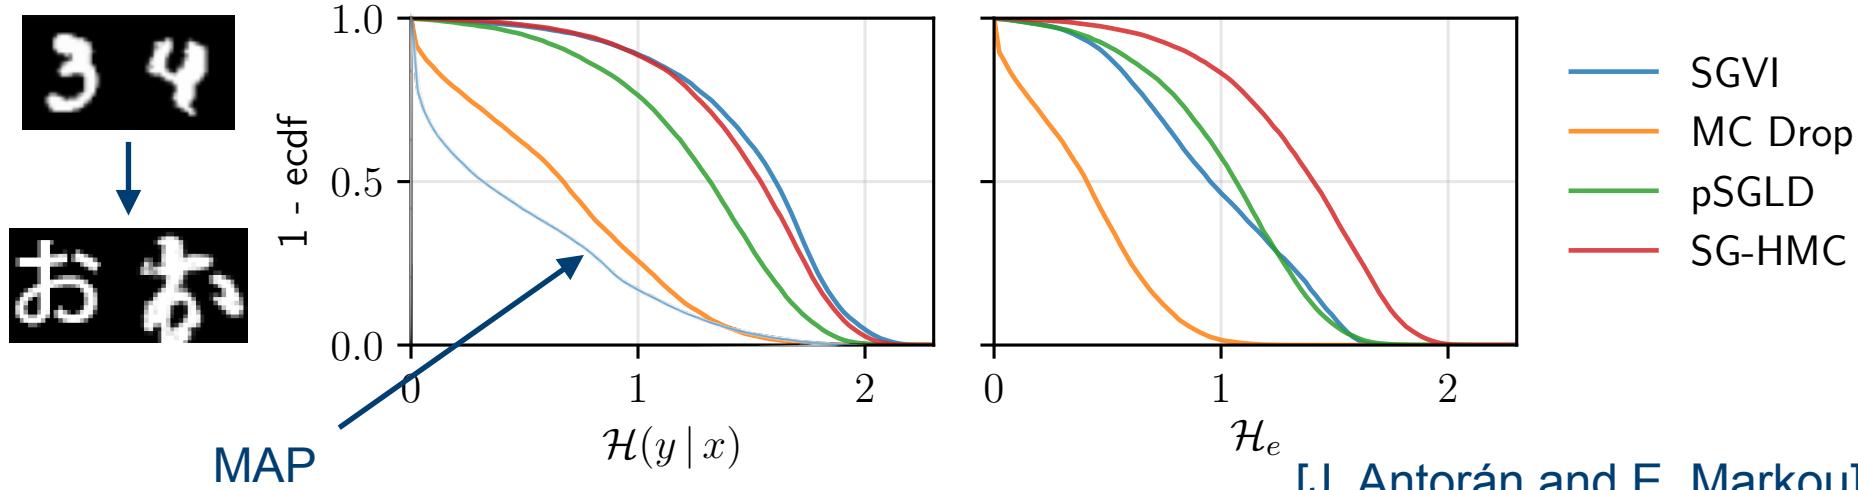
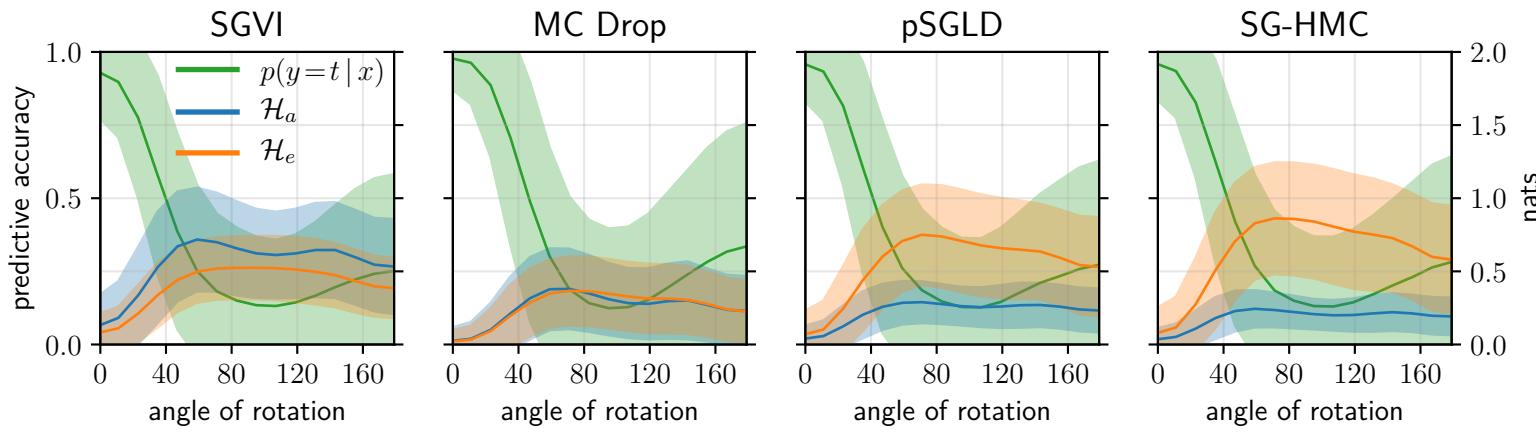
Protection Against Overfitting (cont)

- Plot of train error and test error when training different models with increasing number of permuted samples



[Unpublished results]

Out of Distribution Sample Detection



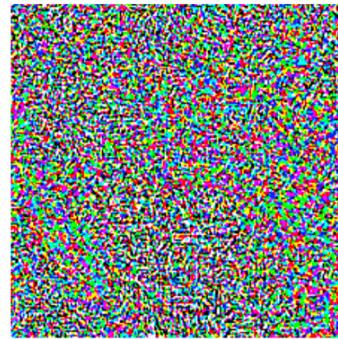
Adversarial Robustness



“panda”

57.7% confidence

$+ .007 \times$



noise

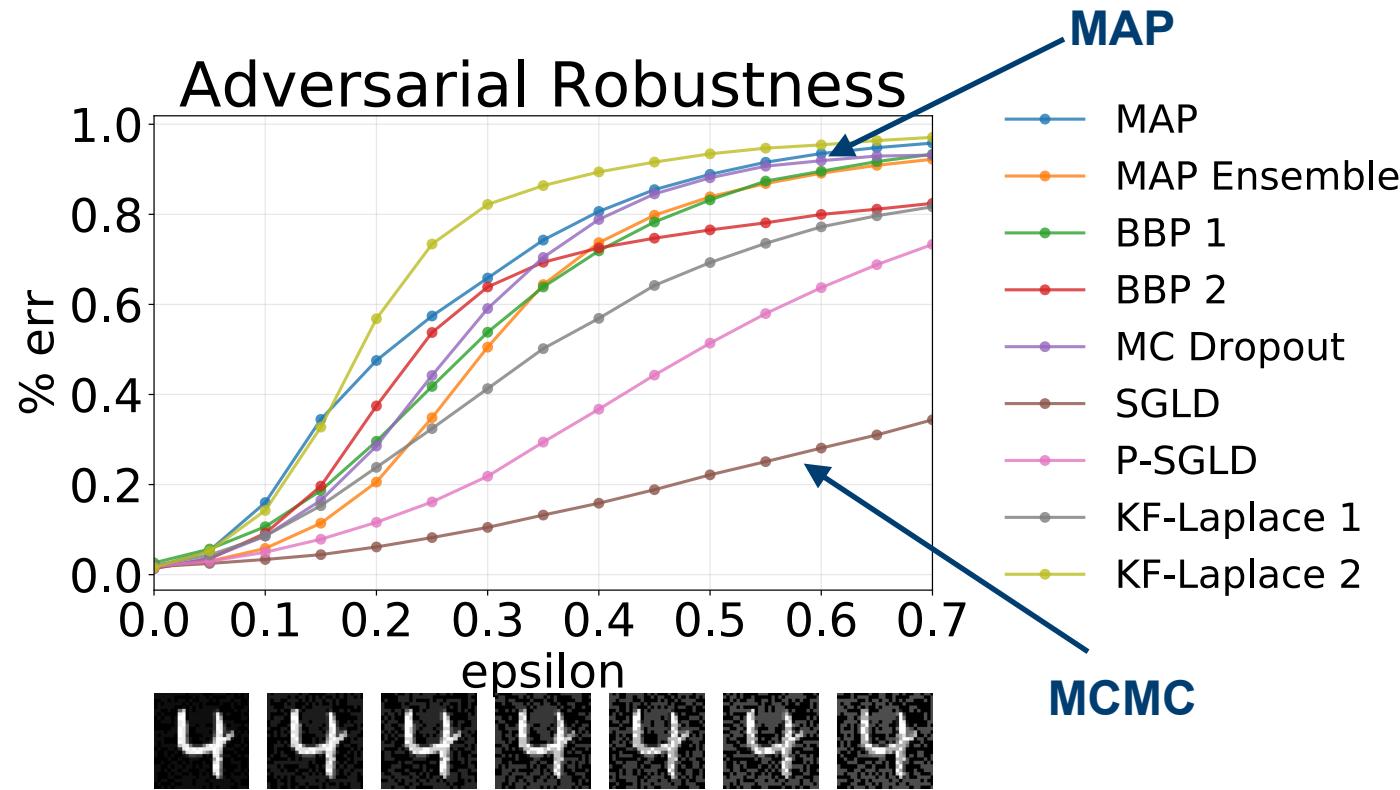
=



“gibbon”

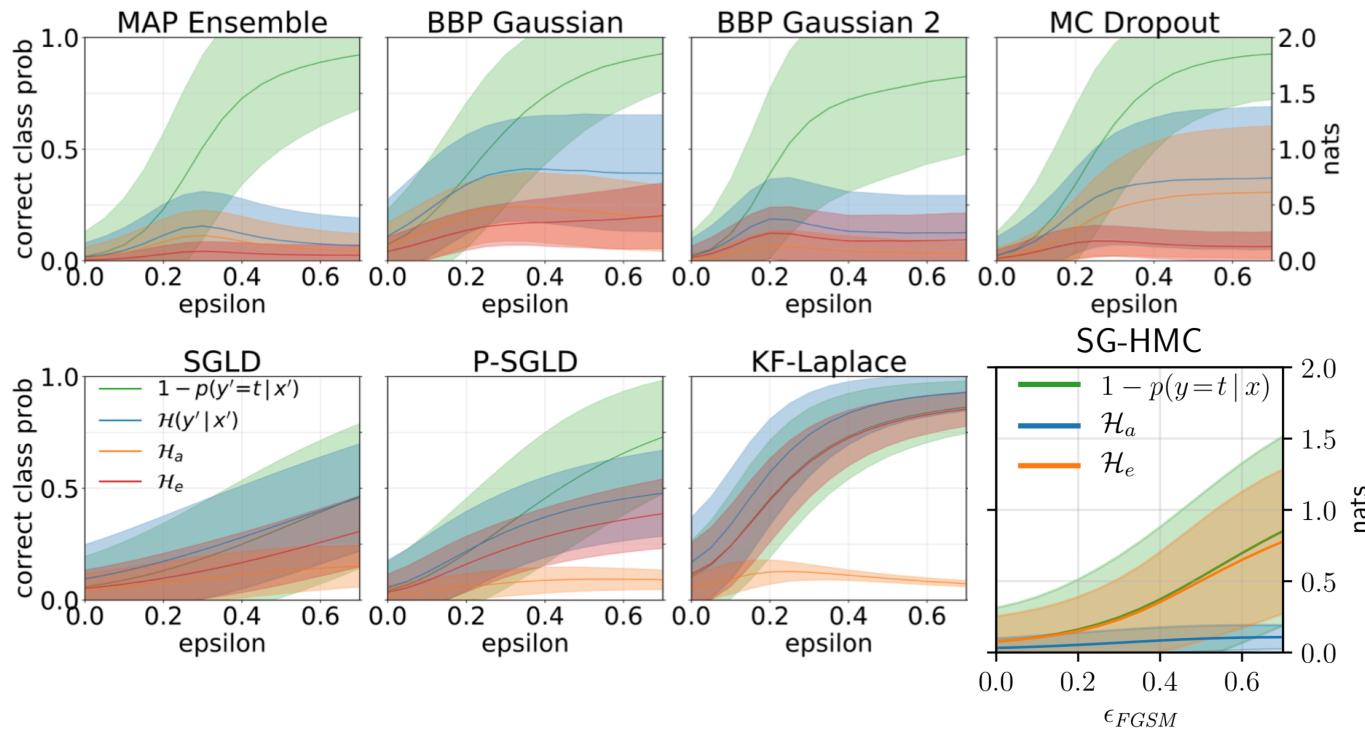
99.3% confidence

Adversarial Robustness



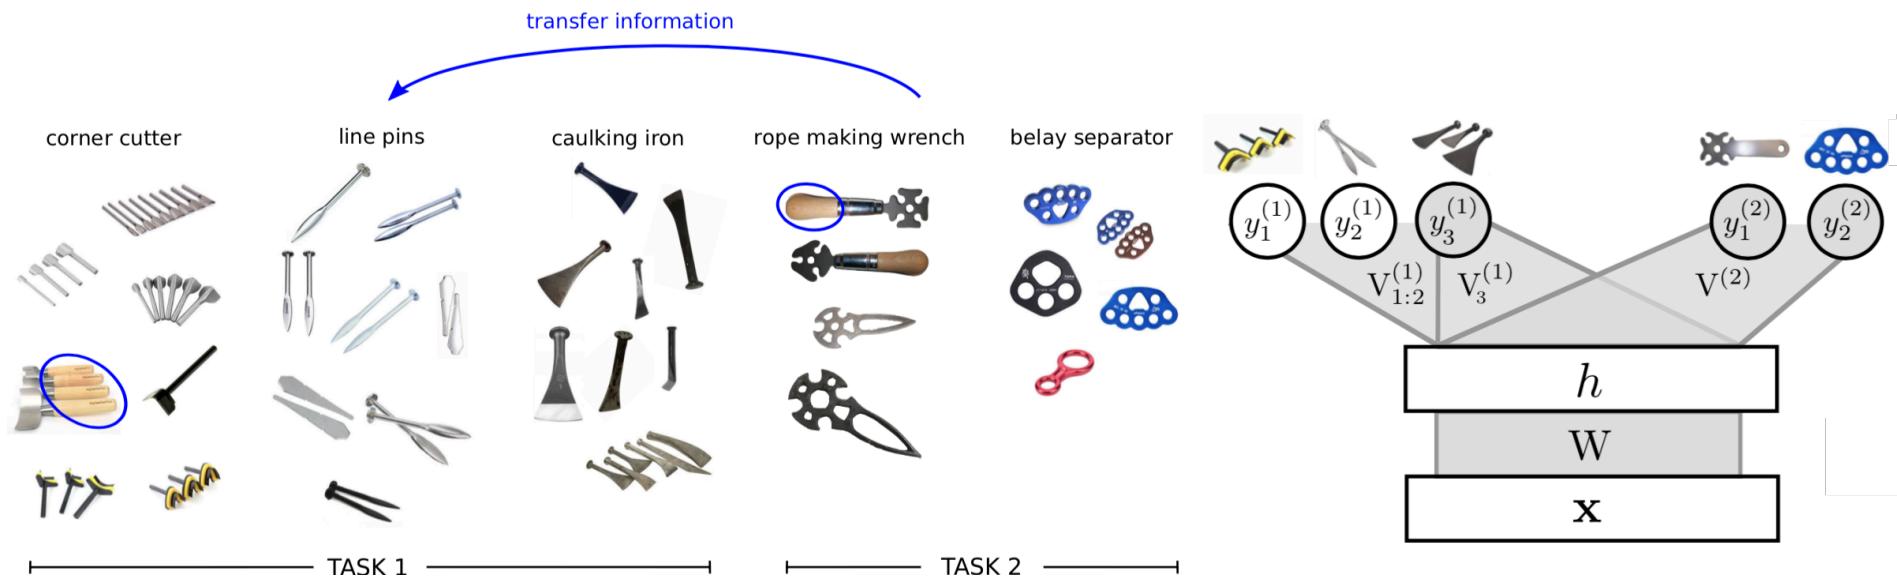
[J. Antorán and E. Markou]

Adversarial Robustness

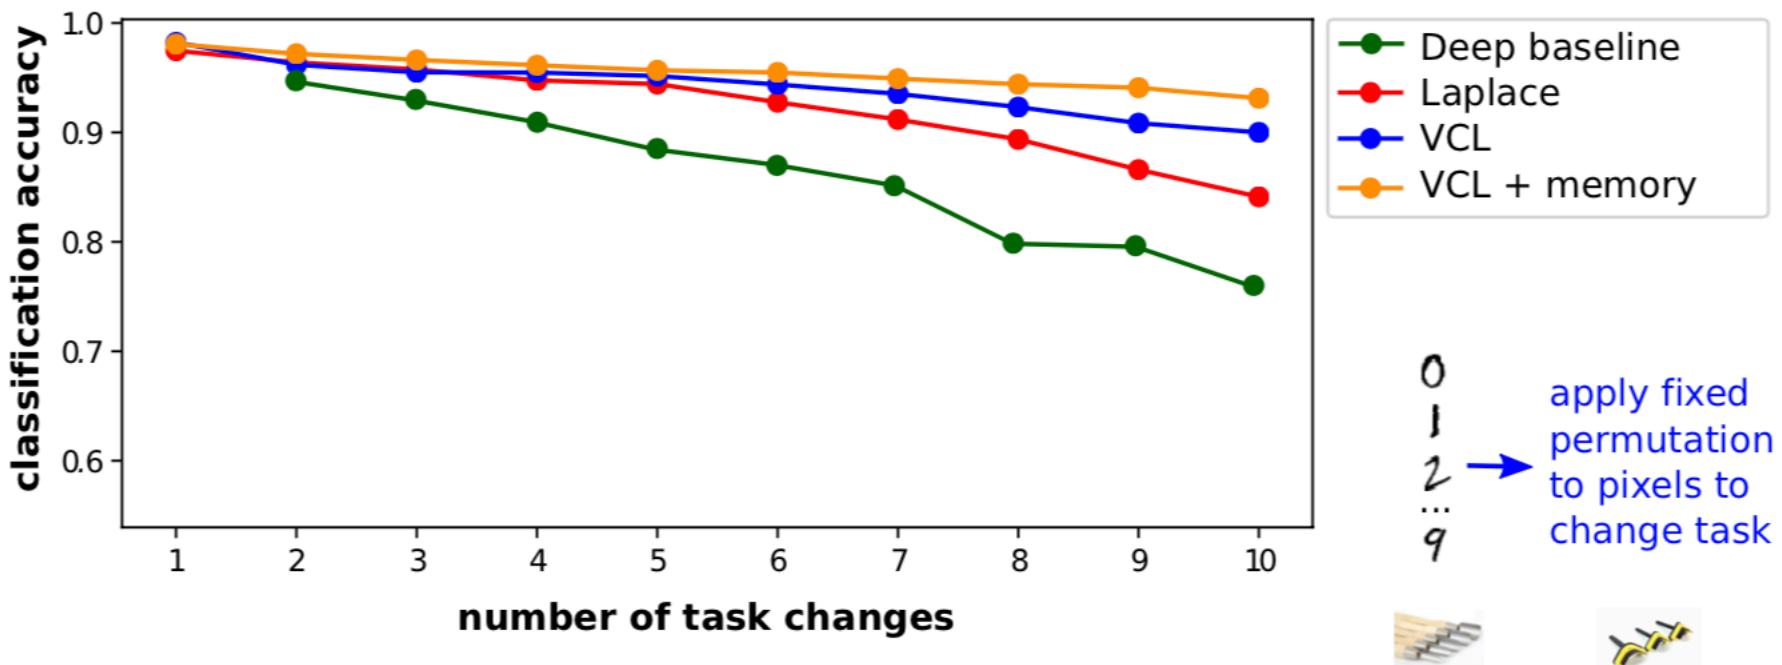


[J. Antorán and E. Markou]

Continual Learning



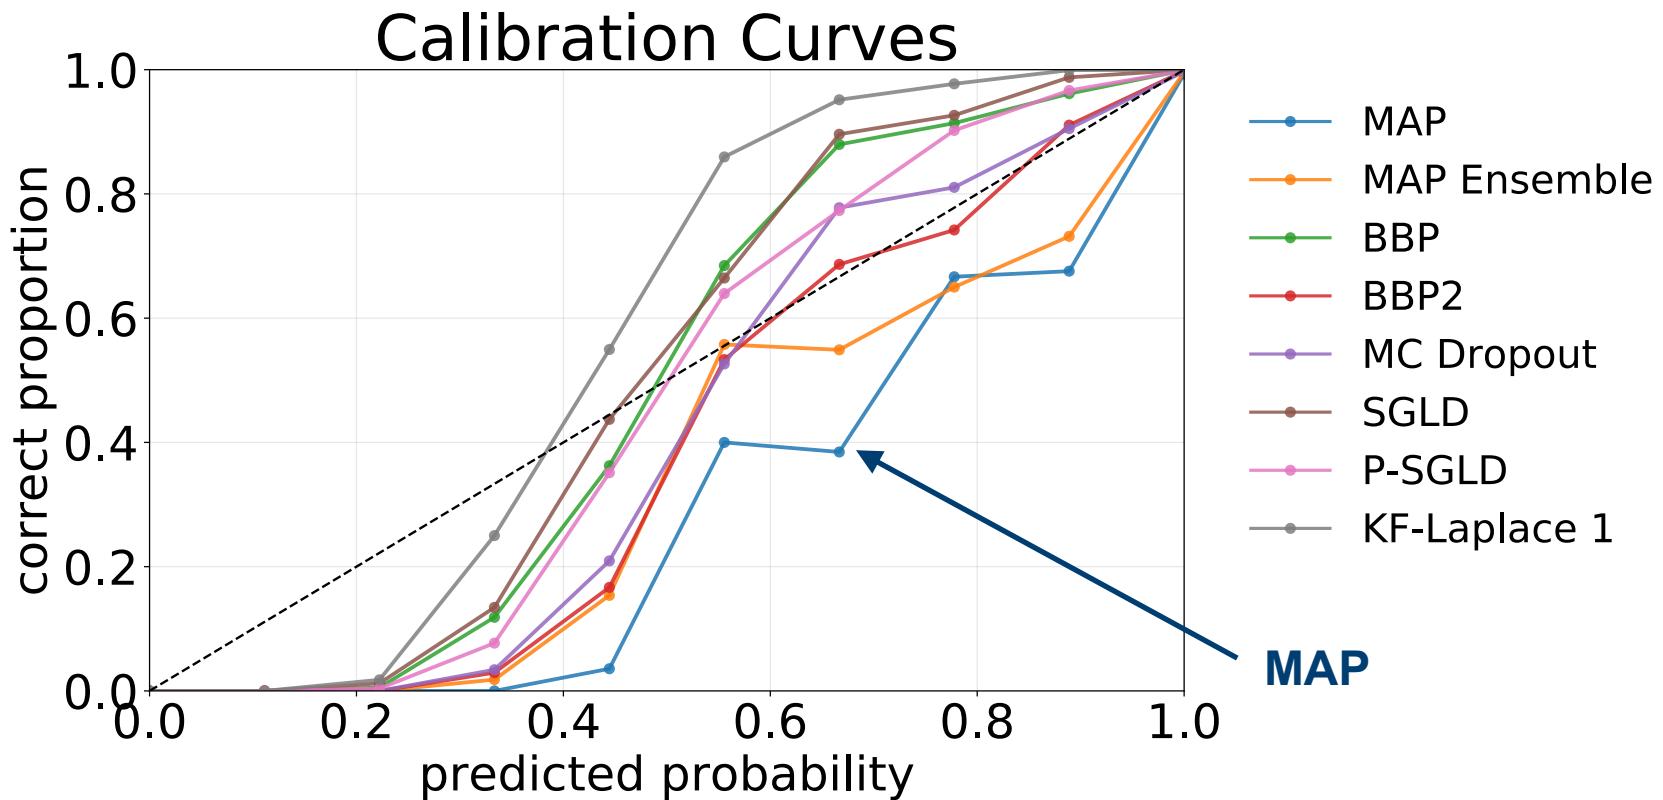
Continual Learning



$$p(\mathbf{w} | \mathcal{D}_1, \mathcal{D}_2) = \frac{p(\mathcal{D}_2 | \mathbf{w})q(\mathbf{w} | \mathcal{D}_1)}{p(\mathcal{D}_2)}$$

[C. Nguyen et al.]

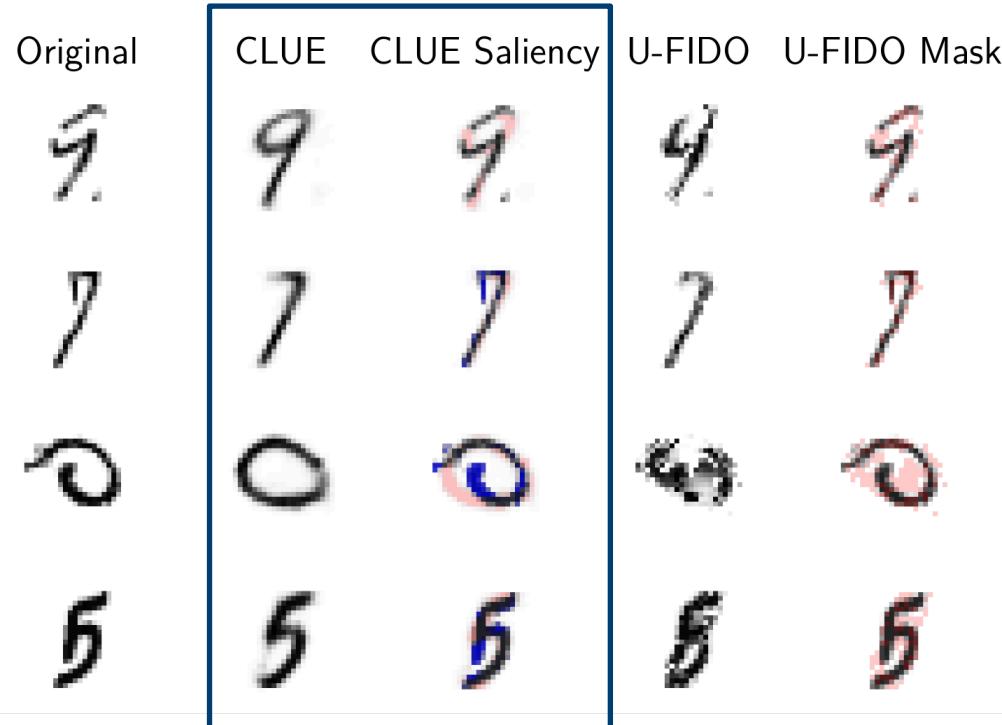
Calibration



[J. Antorán and E. Markou]

Uncertainty Interpretability

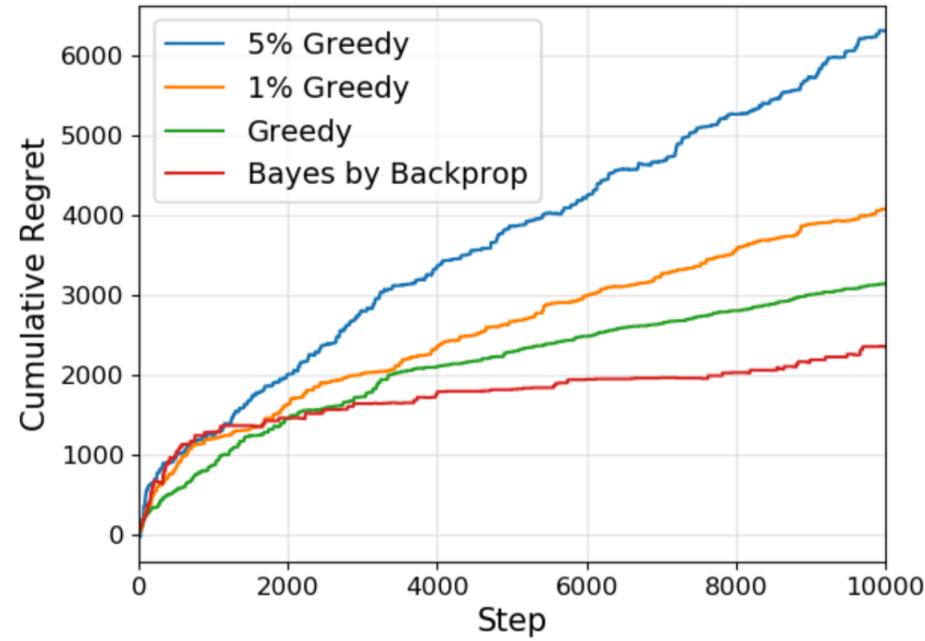
- Why is our model uncertain?



[J. Antorán]

Balancing Exploration and Exploitation: Bandits

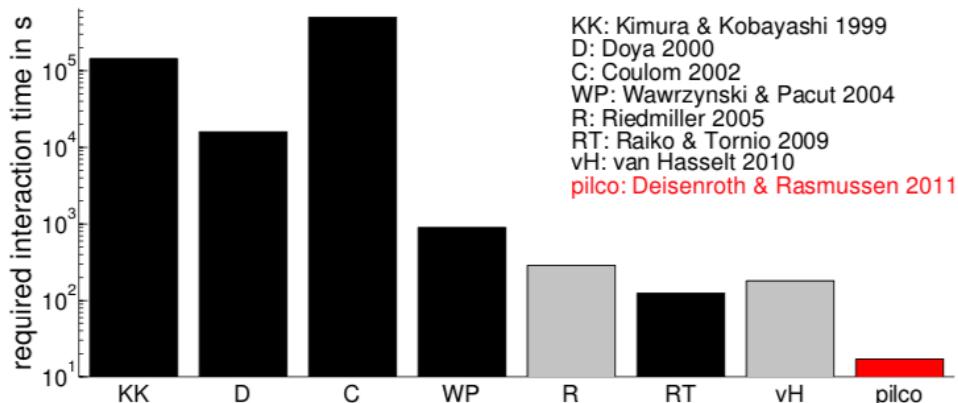
- Different mushrooms give different rewards with different probs.



- + RL / Active Learning / Bayesian Optimisation

Sample-Efficiency in Reinforcement Learning

- Probabilistic models help reduce model bias in RL
- PILCO uses a GP model and a free-form (linear, RBF, NN) controller.



PILCO learner

PILCO learner CS

[M. P. Deisenroth and C. E. Rasmussen]

Open Source BNNs

github.com/JavierAntoran/Bayesian-Neural-Networks

 [Bayesian-Neural-Networks](#) 

Pytorch implementations of Bayes By Backprop, MC Dropout, SGLD, the Local Reparametrization Trick, KF-Laplace and more

 Jupyter Notebook  140  19

Bayesian Neural Networks

License  MIT python 2.7+ pytorch 1.0.1

Pytorch implementations for the following approximate inference methods:

- [Bayes by Backprop](#)
- [Bayes by Backprop + Local Reparametrisation Trick](#)
- [MC dropout](#)
- [Stochastic Gradient Langevin Dynamics](#)
- [Preconditioned SGLD](#)
- [Kronecker-Factorised Laplace Approximation](#)
- Stochastic Gradient Hamiltonian Monte Carlo (Coming soon)

We also provide code for:

- [Bootstrap MAP Ensemble](#)



References

- Some slides taken from Richard E. Turner

Uncertainty in Bayesian Neural Networks

J. Antorán and E. Markou.

Presented at: Workshop on The Mathematics of Deep Learning and Data Science, The Isaac Newton Institute for Mathematical Sciences, Cambridge. 2019.

Excellent textbook introductions to Bayesian machine learning:

- D. J. C. MacKay [Information Theory, Inference and Learning Algorithms](#), 2003
- C. Bishop [Pattern Recognition and Machine Learning](#), 2006
- K. Murphy [Machine Learning: A Probabilistic Perspective](#), 2012

Beautiful seminal work on Bayesian neural networks:

- R. M. Neal [Bayesian Learning for Neural Networks](#), Lecture Notes in Statistics No. 118.
New York: Springer-Verlag. 1996

More References (VI and Laplace)

- D. J. C. MacKay [A practical Bayesian framework for backpropagation networks](#), Neural Computation 4(3) 448-472, 1992
- D. J. C. MacKay [Probable networks and plausible predictions – a review of practical Bayesian methods for supervised neural networks](#), Network: Computation in Neural Systems 6(3), 469-505, 1995
- H. Ritter et al. [A Scalable Laplace Approximation for Neural Networks](#), ICLR, 2018
- G. Hinton and D. Van Camp [Keeping the neural networks simple by minimizing the description length of the weights](#), COLT 5-13, ACM, 1993
- D. Barber and C. Bishop [Ensemble Learning in Bayesian Neural Networks](#), Neural Networks and Machine Learning 1998
- A. Graves [Practical Variational Inference for Neural Networks](#), NIPS 2011
- C. Blundell et al. [Weight Uncertainty in Neural Networks](#), ICML, 2015

More References (HMC)

Bayesian Learning via Stochastic Gradient Langevin Dynamics, Welling and Teh, ICML, 2011

Stochastic Gradient Hamiltonian Monte Carlo, Chen, Fox, Guestrin, NIPS, 2015